

In [ ]: *# Importing the libraries.*

```
import datetime
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, Lasso
from sklearn.metrics import r2_score
from sklearn.impute import SimpleImputer
```

In [ ]: *# Importing Dataset as dataframe.*

```
df = pd.read_csv("dataset.csv")
df.head()
```

Out[ ]:      Unnamed: 0      Name      Location      Year      Kilometers\_Driven      Fuel\_Type      Transmission      C

0	0	Maruti Wagon R LXI CNG	Mumbai	2010	72000	CNG	Manual
1	1	Hyundai Creta 1.6 CRDi SX Option	Pune	2015	41000	Diesel	Manual
2	2	Honda Jazz V	Chennai	2011	46000	Petrol	Manual
3	3	Maruti Ertiga VDI	Chennai	2012	87000	Diesel	Manual
4	4	Audi A4 New 2.0 TDI Multitronic	Coimbatore	2013	40670	Diesel	Automatic



In [ ]: *# Exploring data.*

```
df.describe()
```

Out[ ]:

	Unnamed: 0	Year	Kilometers_Driven	Mileage	Engine	
<b>count</b>	6019.000000	6019.000000	6.019000e+03	6017.000000	5983.000000	5977.00
<b>mean</b>	3009.000000	2013.358199	5.873838e+04	18.134961	1621.276450	5.2
<b>std</b>	1737.679967	3.269742	9.126884e+04	4.582289	601.355233	0.80
<b>min</b>	0.000000	1998.000000	1.710000e+02	0.000000	72.000000	0.00
<b>25%</b>	1504.500000	2011.000000	3.400000e+04	15.170000	1198.000000	5.00
<b>50%</b>	3009.000000	2014.000000	5.300000e+04	18.150000	1493.000000	5.00
<b>75%</b>	4513.500000	2016.000000	7.300000e+04	21.100000	1984.000000	5.00
<b>max</b>	6018.000000	2019.000000	6.500000e+06	33.540000	5998.000000	10.00

In [ ]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6019 entries, 0 to 6018
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            6019 non-null  int64
1   Name                  6019 non-null  object
2   Location              6019 non-null  object
3   Year                  6019 non-null  int64
4   Kilometers_Driven     6019 non-null  int64
5   Fuel_Type             6019 non-null  object
6   Transmission          6019 non-null  object
7   Owner_Type            6019 non-null  object
8   Mileage               6017 non-null  float64
9   Engine                5983 non-null  float64
10  Power                 5983 non-null  object
11  Seats                 5977 non-null  float64
12  New_Price             824 non-null   object
13  Price                 6019 non-null  float64
dtypes: float64(4), int64(3), object(7)
memory usage: 658.5+ KB
```

In [ ]: `# Finding missing values``df.isnull().sum()`

```
Out[ ]: Unnamed: 0      0
        Name          0
        Location      0
        Year          0
        Kilometers_Driven 0
        Fuel_Type      0
        Transmission   0
        Owner_Type     0
        Mileage        2
        Engine         36
        Power          36
        Seats          42
        New_Price      5195
        Price          0
        dtype: int64
```

```
In [ ]: # Dropping rows with null values

df.dropna(inplace=True)
```

```
In [ ]: # Finding missing values

df.isnull().sum()
```

```
Out[ ]: Unnamed: 0      0
        Name          0
        Location      0
        Year          0
        Kilometers_Driven 0
        Fuel_Type      0
        Transmission   0
        Owner_Type     0
        Mileage        0
        Engine         0
        Power          0
        Seats          0
        New_Price      0
        Price          0
        dtype: int64
```

```
In [ ]: # Exploring parameters

print(df.Fuel_Type.value_counts())
print(df.Transmission.value_counts())
print(df.Owner_Type.value_counts())
```

```

Fuel_Type
Diesel      443
Petrol      371
CNG          9
Name: count, dtype: int64
Transmission
Manual       512
Automatic    311
Name: count, dtype: int64
Owner_Type
First        765
Second       55
Third         3
Name: count, dtype: int64

```

```

In [ ]: # Encoding dataframe

# For fuel types.
df.replace({'Fuel_Type': {'Petrol': 0, 'Diesel': 1, 'CNG': 2, 'LPG': 3, 'Ele

# For Transmission type
df.replace({'Transmission': {'Manual': 0, 'Automatic': 1}}, inplace=True)

# For Owner Type
df.replace({'Owner_Type': {'First': 0, 'Second': 1, 'Third': 2, 'Fourth & Ab

```

```

/tmp/ipykernel_358/133526643.py:4: FutureWarning: Downcasting behavior in `r
eplace` is deprecated and will be removed in a future version. To retain the
old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in
to the future behavior, set `pd.set_option('future.no_silent_downcasting', T
rue)`
  df.replace({'Fuel_Type': {'Petrol': 0, 'Diesel': 1, 'CNG': 2, 'LPG': 3, 'E
lectric': 4}}, inplace=True)
/tmp/ipykernel_358/133526643.py:7: FutureWarning: Downcasting behavior in `r
eplace` is deprecated and will be removed in a future version. To retain the
old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in
to the future behavior, set `pd.set_option('future.no_silent_downcasting', T
rue)`
  df.replace({'Transmission': {'Manual': 0, 'Automatic': 1}}, inplace=True)
/tmp/ipykernel_358/133526643.py:10: FutureWarning: Downcasting behavior in `
replace` is deprecated and will be removed in a future version. To retain th
e old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-i
n to the future behavior, set `pd.set_option('future.no_silent_downcasting',
True)`
  df.replace({'Owner_Type': {'First': 0, 'Second': 1, 'Third': 2, 'Fourth &
Above': 3}}, inplace=True)

```

```

In [ ]: # Updated dataset

```

```

df.head()

```

Out[ ]:

	Unnamed: 0	Name	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	Ow
2	2	Honda Jazz V	Chennai	2011	46000	0		0
7	7	Toyota Innova Crysta 2.8 GX AT 8S	Mumbai	2016	36000	1		1
10	10	Maruti Ciaz Zeta	Kochi	2018	25692	0		0
15	15	Mitsubishi Pajero Sport 4X4	Delhi	2014	110000	1		0
20	20	BMW 3 Series 320d	Kochi	2014	32982	1		1

In [ ]: *# Splitting the dataframe into data and target.*

```
X = df.drop(['Price', 'Name', 'New_Price', 'Location'], axis=1)
Y = df['Price']
```

In [ ]: X.head()

Out[ ]:

	Unnamed: 0	Year	Kilometers_Driven	Fuel_Type	Transmission	Owner_Type	Mileage	E
2	2	2011	46000	0	0	0	18.20	1
7	7	2016	36000	1	1	0	11.36	2
10	10	2018	25692	0	0	0	21.56	1
15	15	2014	110000	1	0	0	13.50	2
20	20	2014	32982	1	1	0	22.69	1

In [ ]: Y.head()

Out[ ]:

```
2      4.50
7     17.50
10     9.95
15    15.00
20    18.55
Name: Price, dtype: float64
```

In [ ]: *# Splitting data into training and testing sets.*

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2, r
```

```
In [ ]: # Fitting the LinearRegression model
```

```
model = LinearRegression()  
model.fit(X_train, Y_train)
```

```
Out[ ]: ▼ LinearRegression ⓘ ⓘ  
LinearRegression()
```

```
In [ ]: # Implementing Prediction
```

```
prediction = model.predict(X_train)
```


```
In [ ]: # Calculating model's accuracy
```

```
error_score = r2_score(Y_train, prediction)  
  
error_score
```

```
Out[ ]: 0.6863803349632541
```

```
In [ ]: # Representing accuracy visually using scatter plot fro training data
```

```
plt.scatter(Y_train, prediction)  
plt.xlabel("Actual Price")  
plt.ylabel("Predicted Price")  
plt.title("Actual vs Predicted Prices")  
plt.show()
```




```
In [ ]: # Predicting for testing dat  
  
test_pred = model.predict(X_test)
```

```
In [ ]: # Calculating accuracy for model using testing data  
  
error_score = r2_score(Y_test, test_pred)  
  
error_score
```

```
Out[ ]: 0.6997672112263285
```

```
In [ ]: # Representing accuracy visually using scatter plot fro testing data  
  
plt.scatter(Y_test, test_pred)  
plt.xlabel("Actual Price")  
plt.ylabel("Predicted Price")  
plt.title("Actual vs Predicted Prices")  
plt.show()
```



In [ ]: *# Importing new data*

```
nd = pd.read_csv("new_data.csv")
nd.head()
```

Out [ ]:

Unnamed: 0	Name	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	Owner_Type
0	Honda City (V) 2015	Shahjahanpur	2015	89000	Diesel	Manual	First

In [ ]: *# Encoding new data*

```
# For fuel types.
nd.replace({'Fuel_Type': {'Petrol': 0, 'Diesel': 1, 'CNG': 2, 'LPG': 3, 'Electricity': 4}})

# For Transmission type
nd.replace({'Transmission': {'Manual': 0, 'Automatic': 1}}, inplace=True)

# For Owner Type
nd.replace({'Owner_Type': {'First': 0, 'Second': 1, 'Third': 2, 'Fourth & Above': 3}})

nd
```



```

/tmp/ipykernel_358/1013661178.py:4: FutureWarning: Downcasting behavior in `
replace` is deprecated and will be removed in a future version. To retain th
e old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-i
n to the future behavior, set `pd.set_option('future.no_silent_downcasting'
, True)`
    nd.replace({'Fuel_Type': {'Petrol': 0, 'Diesel': 1, 'CNG': 2, 'LPG': 3, 'E
lectric': 4}}, inplace=True)
/tmp/ipykernel_358/1013661178.py:7: FutureWarning: Downcasting behavior in `
replace` is deprecated and will be removed in a future version. To retain th
e old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-i
n to the future behavior, set `pd.set_option('future.no_silent_downcasting'
, True)`
    nd.replace({'Transmission': {'Manual': 0, 'Automatic': 1}}, inplace=True)
/tmp/ipykernel_358/1013661178.py:10: FutureWarning: Downcasting behavior in
`replace` is deprecated and will be removed in a future version. To retain t
he old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-
in to the future behavior, set `pd.set_option('future.no_silent_downcastin
g', True)`
    nd.replace({'Owner_Type': {'First': 0, 'Second': 1, 'Third': 2, 'Fourth &
Above': 3}}, inplace=True)

```

```

Out[ ]:      Unnamed: 0  Name      Location  Year  Kilometers_Driven  Fuel_Type  Transmission  Ow
0      0      Honda City (V) 2015      89000      1      0

```

```

In [ ]: # Pre-processing new data

pred_data = nd.drop(['Name', 'Location', 'Price'], axis = 1)
pred_data

```

```

Out[ ]:      Unnamed: 0  Year  Kilometers_Driven  Fuel_Type  Transmission  Owner_Type  Mileage  En
0      0      2015      89000      1      0      0      21

```

```

In [ ]: # Predicting based on new data

new_predicton = model.predict(pred_data)
new_predicton

```

```

Out[ ]: array([10.82642432])

```

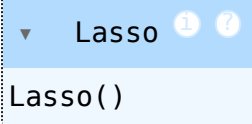
LASSO REGRESSION MODEL

```

In [ ]: # Fitting lasso model

lasso_model = Lasso()
lasso_model.fit(X_train, Y_train)

```

Out[ ]:  Lasso()  
Lasso()

```
In [ ]: # Predicting via lasso model  
  
lasso_pred = lasso_model.predict(X_train)
```

```
In [ ]: # Calculating lasso model's accuracy  
  
error_score = r2_score(Y_train, lasso_pred)  
  
error_score
```

Out[ ]: 0.6646931120946568

```
In [ ]: # Representing lasso model's accuracy visually  
  
plt.scatter(Y_train, lasso_pred)  
plt.xlabel("Actual Price")  
plt.ylabel("Predicted Price")  
plt.title("Actual vs Predicted Prices (Lasso Model)")  
plt.show()
```

